

# LLVM Sanitizers & GDB(gui)

***Nils Wentzell***

*Center for Computational Quantum Physics (CCQ)*

*Flatiron Institute, Simons Foundation*

*New York*

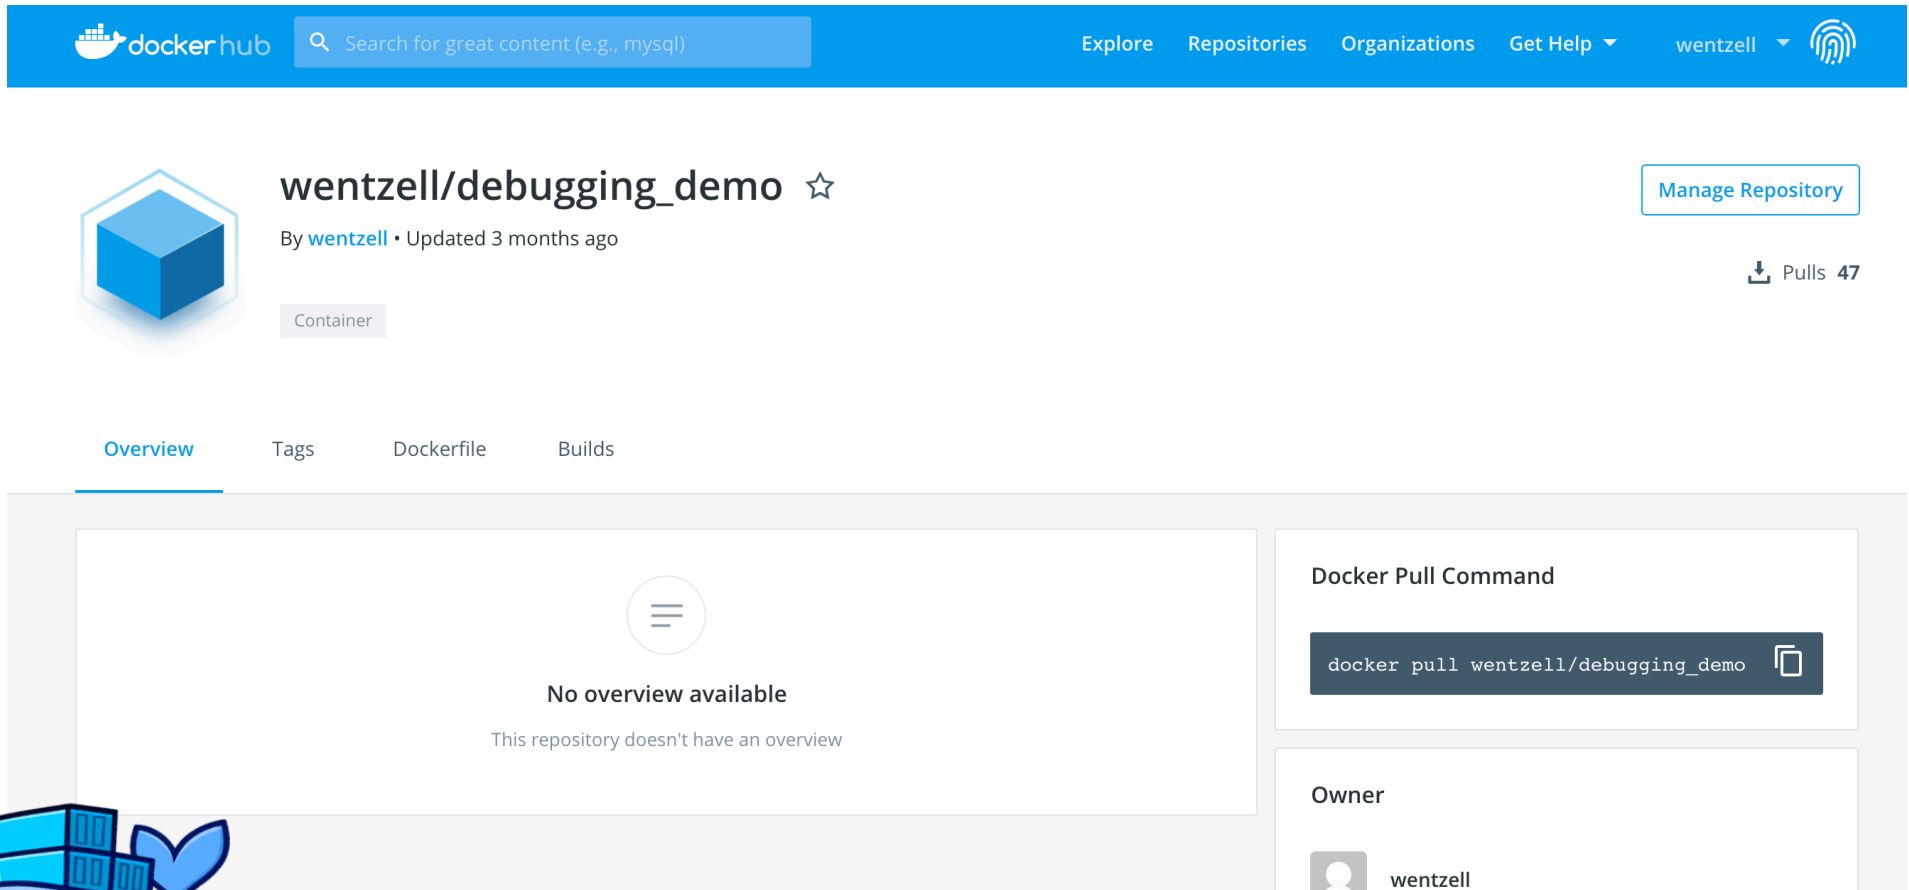
SIMONS FOUNDATION



# Docker Image

[hub.docker.com/r/wentzell/debugging\\_demo](https://hub.docker.com/r/wentzell/debugging_demo)

[github.com/wentzell/debugging\\_demo](https://github.com/wentzell/debugging_demo)



The screenshot shows the Docker Hub interface for the repository `wentzell/debugging_demo`. The page has a blue header with the Docker Hub logo, a search bar, and navigation links: Explore, Repositories, Organizations, Get Help, and a user profile for `wentzell`. The repository card features a blue cube icon, the name `wentzell/debugging_demo` with a star icon, and text indicating it was updated 3 months ago. A 'Container' tag is visible. A 'Manage Repository' button is in the top right. Below the card, tabs for Overview, Tags, Dockerfile, and Builds are shown. The main content area displays 'No overview available' with a message that the repository doesn't have an overview. On the right, the 'Docker Pull Command' is shown as `docker pull wentzell/debugging_demo`, and the 'Owner' is listed as `wentzell`. A pull count of 47 is also shown.

**`docker pull wentzell/debugging_demo`**

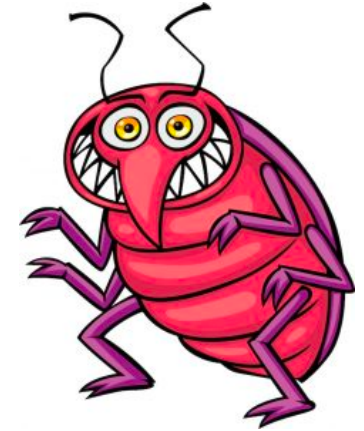


# Avoiding bugs

- Expressive Code (Modern C++)
- Code Review
- Automated Tests (googletest, TDD)
- Version Control (git)
- Static analyzer's (clang-tidy)



# Avoiding bugs



- Expressive Code (Modern C++)
  - Code Review
  - Automated Tests (googletest, TDD)
  - Version Control (git)
  - Static analyzer's (clang-tidy)
- Use dynamic analyzer tools to catch them as they appear!
    - Valgrind
    - Clang Sanitizers

- Set of dynamic analysis tools:
  - Memcheck (default) `--tool=memcheck`  
Memory error detector  
(Illegal read/writes, memory leaks,  
use of uninitialized memory, ... )
  - Cachegrind: Cache profiler `--tool=cachegrind`
  - Callgrind: Callgraph analyzer `--tool=callgrind`
  - Massif: Heap profiler `--tool=massif`



- Set of dynamic analysis tools:

- Memcheck (default) `--tool=memcheck`  
Memory error detector  
(Illegal read/writes, memory leaks,  
use of uninitialized memory, ... )

- Cachegrind: Cache profiler `--tool=cachegrind`
- Callgrind: Callgraph analyzer `--tool=callgrind`
- Massif: Heap profiler `--tool=massif`



# Clang sanitizers



[clang.llvm.org](http://clang.llvm.org)

- Dynamic analysis tools included with clang and gcc

- Address Sanitizer `-fsanitize=address`  
detects buffer overflows, memory leaks, use-after-free,...
- Memory Sanitizer (clang only) `-fsanitize=memory`  
detects reads of uninitialized memory
- Undefined Behavior Sanitizer `-fsanitize=undefined`  
detects undefined behaviors
- Thread Sanitizer `-fsanitize=thread`  
detects data races

# Clang sanitizers



[clang.llvm.org](http://clang.llvm.org)

- Dynamic analysis tools included with clang and gcc
- **Address Sanitizer** `-fsanitize=address`  
detects buffer overflows, memory leaks, use-after-free,...
  - **Memory Sanitizer (clang only)** `-fsanitize=memory`  
detects reads of uninitialized memory
  - **Undefined Behavior Sanitizer** `-fsanitize=undefined`  
detects undefined behaviors
  - **Thread Sanitizer** `-fsanitize=thread`  
detects data races
- Used to systematically detect bugs by Google, Mozilla, ...
  - Much less memory/runtime overhead than Valgrind!



# Comparison

## ASan/MSan vs Valgrind (Memcheck)

	Valgrind	ASan	MSan
Heap out-of-bounds	YES	YES	
Stack out-of-bounds		YES	
Global out-of-bounds		YES	
Use-after-free	YES	YES	
Use-after-return		YES	
Uninitialized reads	YES		YES
CPU Overhead	10x-300x	1.5x-3x	3x

Source: Cppcon 2014, Kostya Serebryany

# Examples

# I-Out of bounds !

- Typical error : read/write outside of the memory of an object.

```
#include <vector>

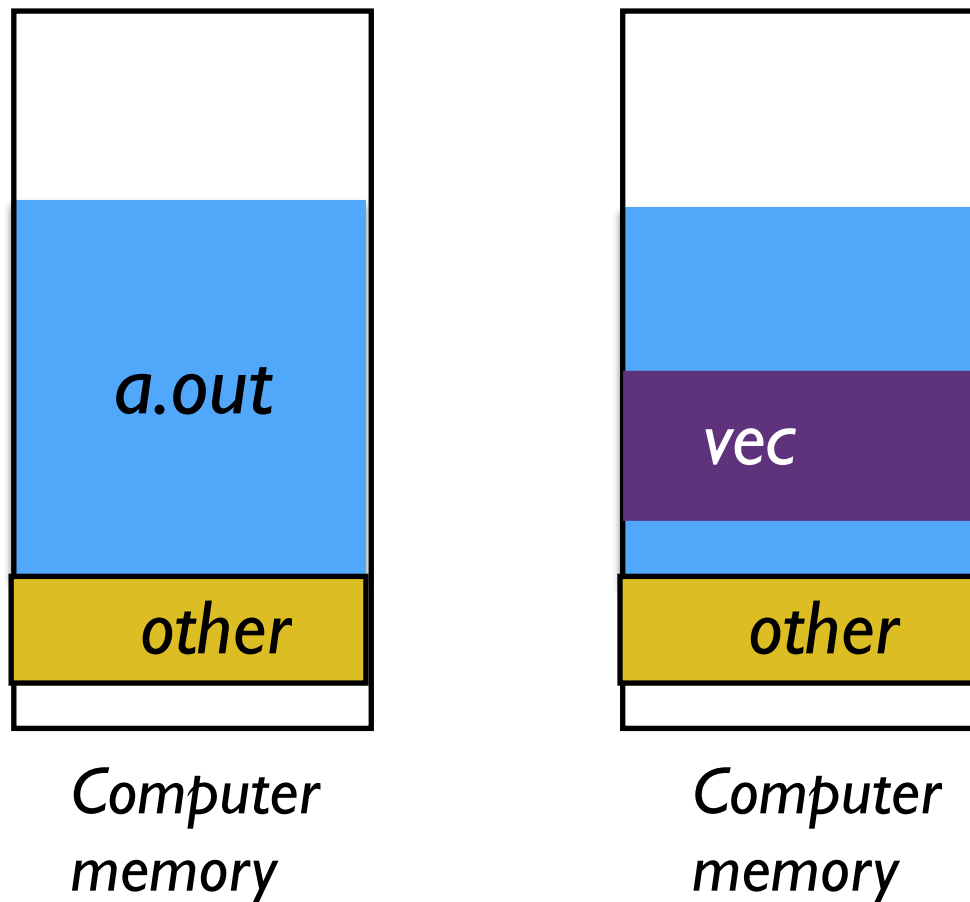
int fun(int n) {
    std::vector<int> vec(100, 1); // a 1d "array" of 100 int init to 1
    vec[n + 50] = 500;           // Invalid Write
    return vec[n + 50];          // Invalid Read
}
```

- No bound-checking by default!
- Similar : use a “dangling” or null pointer, ...
- How does it manifest itself ?

Demo

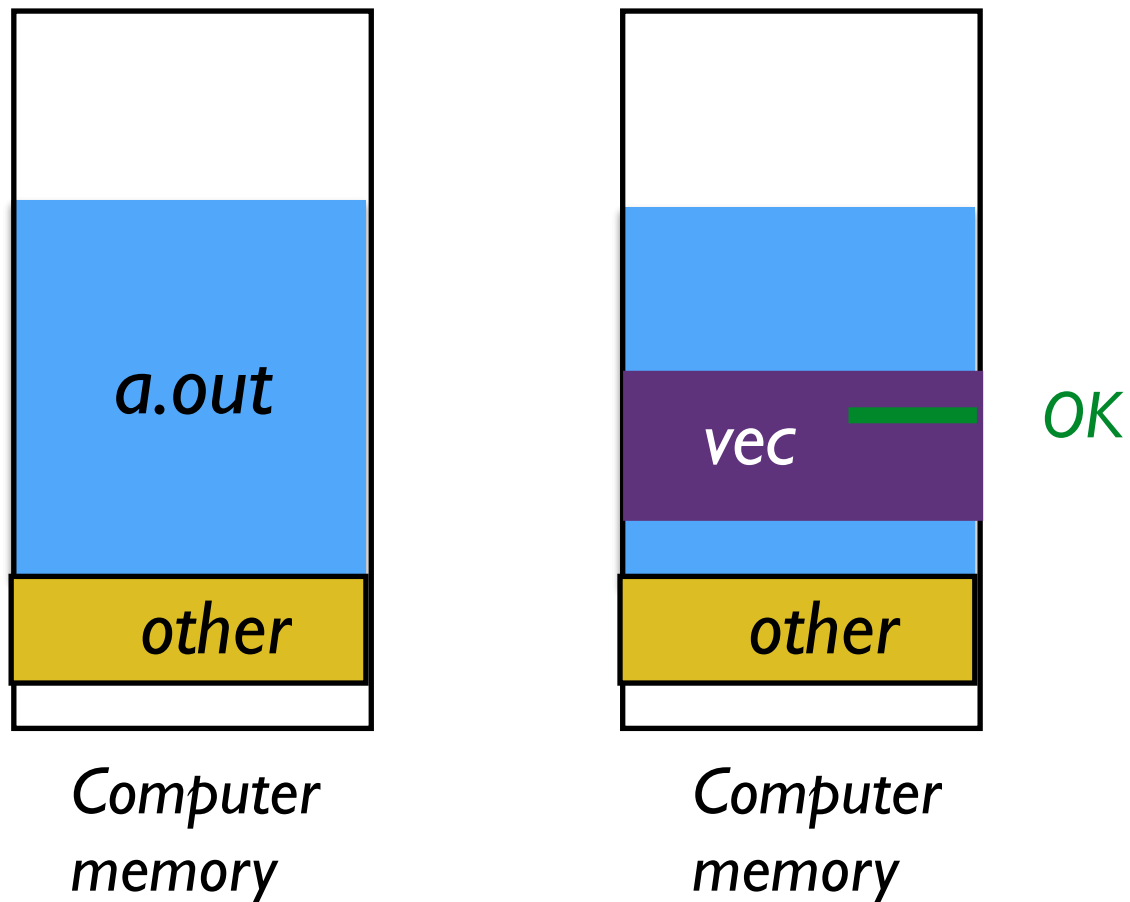
# I-Out of bounds !

- Segfault : the program crashes ... or not !?
- Error can be silent!
- We need to find the first error and stop



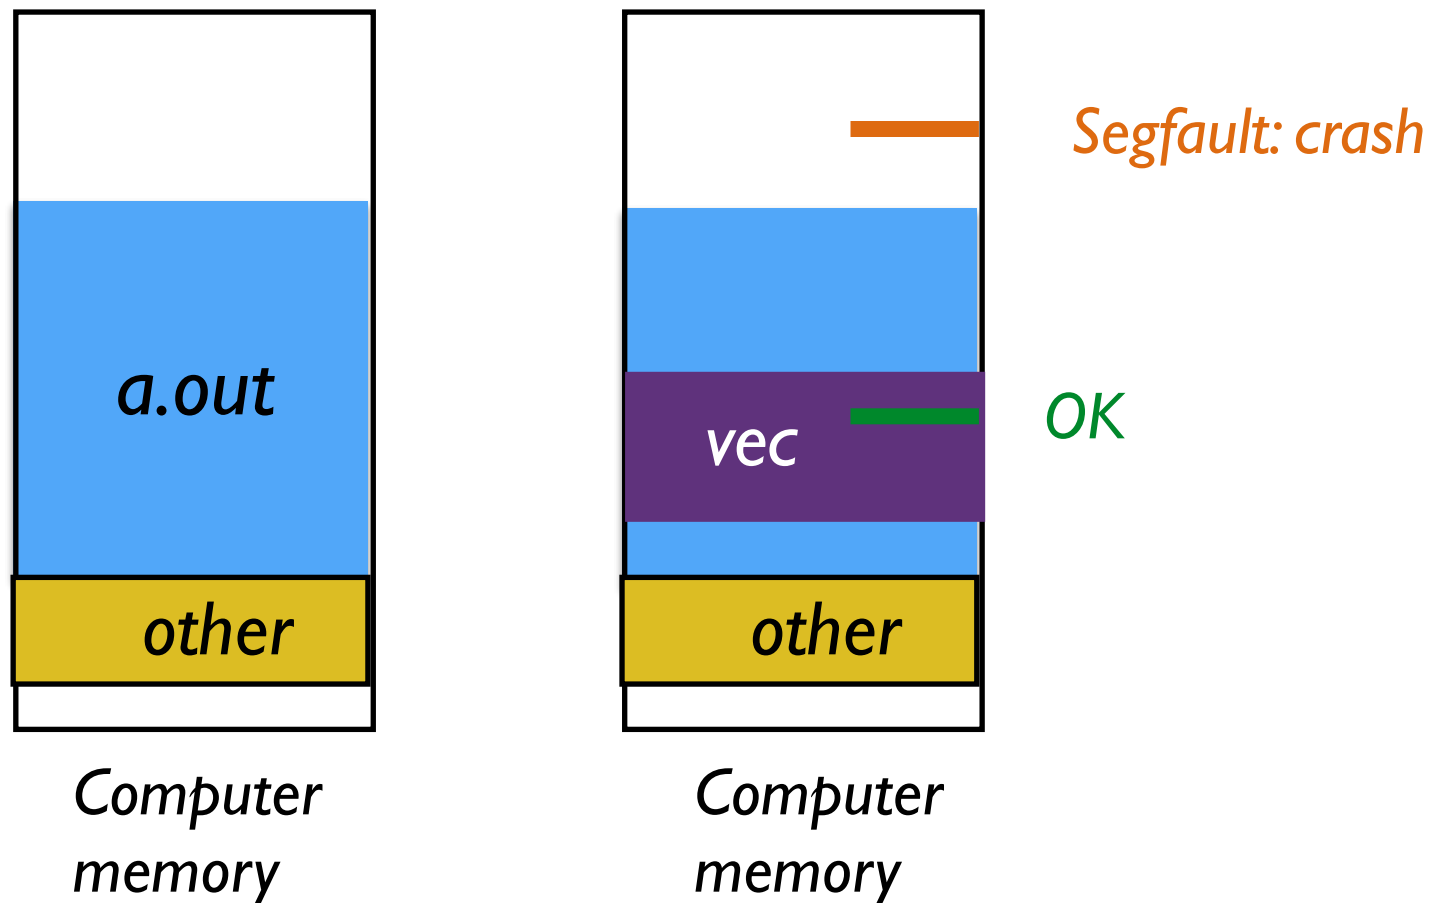
# I-Out of bounds !

- Segfault : the program crashes ... or not !?
- Error can be silent!
- We need to find **the first error** and **stop**



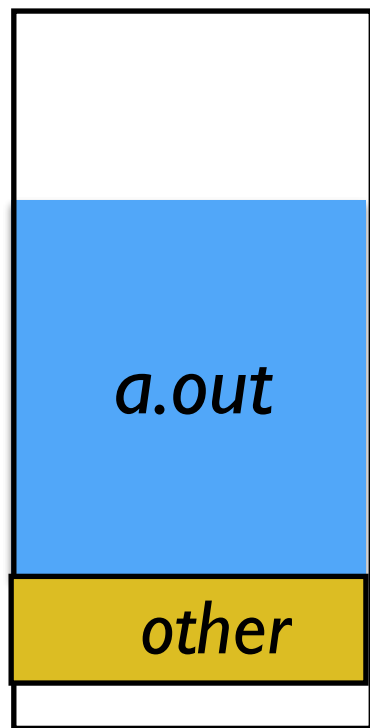
# I-Out of bounds !

- Segfault : the program crashes ... or not !?
- Error can be silent!
- We need to find **the first error** and **stop**

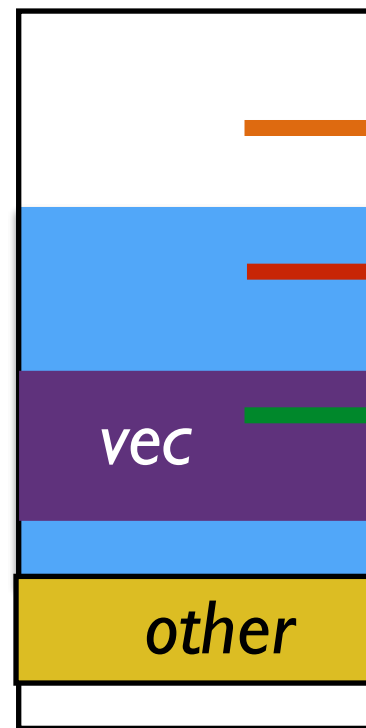


# I-Out of bounds !

- Segfault : the program crashes ... or not !?
- Error can be silent!
- We need to find the first error and stop



Computer  
memory



Computer  
memory

*Segfault: crash*

*Even worse. No crash,  
but corrupt memory*

*OK*

## 2-Invalid-iterator

- A Typical error in C++. Pointer or iterator invalidation.

```
#include <iostream>
#include <vector>

int main() {
    std::vector<int> v(10, 3);           // 10 numbers init to 3
    auto p = v.begin();                 // an iterator on the beginning of the vector
    std::cout << *p << std::endl;      // print the first value : should be 3
    v.push_back(7);                     // Add one more number to the vector,
    std::cout << *p << std::endl;      // print again the first value ???
}
```



## 2-Invalid-iterator

- A Typical error in C++. Pointer or iterator invalidation.

```
#include <iostream>
#include <vector>

int main() {
    std::vector<int> v(10, 3);    // 10 numbers init to 3
    auto p = v.begin();          // an iterator on the beginning of the vector
    std::cout << *p << std::endl; // print the first value : should be 3
    v.push_back(7);              // Add one more number to the vector,
    std::cout << *p << std::endl; // print again the first value ???
}
```

- Next generation of compiler will detect this.  
Lifetime Proposal for std C++.
- Experimental branch of clang.

→ clang++ -Wlifetime 7-IteratorInvalid.cpp

<https://godbolt.org/z/Z878Dp>

# 3- Memory leak

- Allocate some memory, and ... forget to release it

```
int main() {  
    int *g = new int[100]; // allocate the array of 100 int on the heap  
    g = nullptr;           // Lost the pointer.  
  
    // ... whatever ...  
}
```

- **Effect** : Often none, except if you call such a function a lot, you will run out of memory !
- **Modern C++**: no new/delete. Should not pass code review !
- **Tools** : Valgrind, ASAN (address sanitizer).
  - `valgrind --leak-check=full ./a.out`
  - `clang++ -fsanitize=address -g Leak.cpp`

## 4-Uninitialized variables

- What happens if we forgot to initialize something ?

```
#include <iostream>

int main(int argc, char** argv) {
    int num = atoi(argv[1]); // get the first arg and make it from string -> int
    int factorial; // OOPS !
    for (int i = 1; i <= num; ++i) {
        factorial *= i;
    }
    std::cout << num << "! = " << factorial << "\n";
    return factorial;
}
```

- Static analyser
  - Compiler (clang -Wall detects it, but not gcc)
  - Cppcheck
- Dynamical analyzer: Valgrind, memory sanitiser (MSAN)
- **Limitation** : libraries must be compiled with `-fsanitize=memory`

Demo

→ `clang++ -g -fsanitize=memory -fno-omit-frame-pointer -fsanitize-memory-track-origins factorial.cpp`

## 5- Undefined behaviour

- Undefined behavior sanitizer detects various problems:
  - Overflow
  - Division by zero
  - Wrong cast e.g. in calling C lib
  - ... Undefined Behavior situation in C++...

## 5- Undefined behavior

- Integer overflow.

```
#include <iostream>

int main(int argc, char** argv) {
    int num = atoi(argv[1]); // get the first arg and make it from string -> int

    int t = num << 16;
    int r = t * t;
    std::cout << r << std::endl;

    double x = 1/0.0;
    std::cout << x << std::endl;
}
```

```
→ clang++ -Wall int_overflow.cpp
→ ./a.out 5
0
→ clang++ -Wall -fsanitize=undefined int_overflow.cpp
→ ./a.out 5
int_overflow.cpp:7:18:
runtime error: signed integer overflow: 327680 * 327680 cannot be represented in type 'int'
0
```

## 6-Logic error

- The logic of the code is flawed, but no automatic tool will find this !

```
#include <iostream>

int main(int argc, char** argv) {
    int num = atoi(argv[1]); // get the first arg and make it from string -> int

    int factorial = 0; // Oops !
    for (int i = 1; i <= num; ++i) {
        factorial *= i;
    }
    std::cout << num << "! = " << factorial << "\n";
    return factorial;
}
```

Demo

- Let's follow step by step and use the debugger.
- GDB : for many compiled languages: C, C++, FORTRAN, Java, Pascal, Ada, D, Go
- GDBgui : a light browser-based frontend to GDB [www.gdbgui.com](http://www.gdbgui.com)

# 7-Infinite Loop

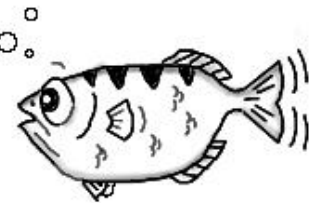
- Due to some logic flaw, your code is stuck in an infinite loop !

```
int main() {  
  
    int c = 0;  
  
    while (1) {  
        c += 1;  
        // whatever  
    }  
}
```

- You can run, and attach gdb to your process on the fly

```
→ gdb ./a.out --pid=PID_OF_THE_PROCESS
```

*Demo*



# GDB: The GNU Project Debugger

18

## GDB cheatsheet - page 1

### Running

```
# gdb <program> [core dump]
    Start GDB (with optional core dump).

# gdb --args <program> <args...>
    Start GDB and pass arguments

# gdb --pid <pid>
    Start GDB and attach to process.

set args <args...>
    Set arguments to pass to program to
    be debugged.

run
    Run the program to be debugged.

kill
    Kill the running program.
```

### Breakpoints

```
break <where>
    Set a new breakpoint.

delete <breakpoint#>
    Remove a breakpoint.

clear
    Delete all breakpoints.

enable <breakpoint#>
    Enable a disabled breakpoint.

disable <breakpoint#>
    Disable a breakpoint.
```

### Watchpoints

```
watch <where>
    Set a new watchpoint.

delete/enable/disable <watchpoint#>
    Like breakpoints.
```

### <where>

```
function_name
    Break/watch the named function.

line_number
    Break/watch the line number in the cur-
    rent source file.

file:line_number
    Break/watch the line number in the
    named source file.
```

### Conditions

```
break/watch <where> if <condition>
    Break/watch at the given location if the
    condition is met.
    Conditions may be almost any C ex-
    pression that evaluate to true or false.

condition <breakpoint#> <condition>
    Set/change the condition of an existing
    break- or watchpoint.
```

### Examining the stack

```
backtrace
where
    Show call stack.

backtrace full
where full
    Show call stack, also print the local va-
    riables in each frame.

frame <frame#>
    Select the stack frame to operate on.
```

### Stepping

```
step
    Go to next instruction (source line), di-
    ving into function.
```

```
next
    Go to next instruction (source line) but
    don't dive into functions.

finish
    Continue until the current function re-
    turns.

continue
    Continue normal execution.
```

### Variables and memory

```
print/format <what>
    Print content of variable/memory locati-
    on/register.

display/format <what>
    Like „print“, but print the information
    after each stepping instruction.

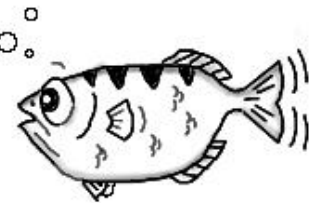
undisplay <display#>
    Remove the „display“ with the given
    number.

enable display <display#>
disable display <display#>
    En- or disable the „display“ with the gi-
    ven number.

x/nfu <address>
    Print memory.
    n: How many units to print (default 1).
    f: Format character (like „print“).
    u: Unit.

    Unit is one of:
    b: Byte,
    h: Half-word (two bytes)
    w: Word (four bytes)
    g: Giant word (eight bytes)).
```





# GDB: The GNU Project Debugger

19

## GDB cheatsheet - page 2

### Format

<i>a</i>	<b>Pointer.</b>
<i>c</i>	Read as integer, print as character.
<i>d</i>	Integer, signed decimal.
<i>f</i>	Floating point number.
<i>o</i>	Integer, print as octal.
<i>s</i>	Try to treat as C string.
<i>t</i>	Integer, print as binary ( <i>t</i> = „two“).
<i>u</i>	Integer, unsigned decimal.
<i>x</i>	Integer, print as hexadecimal.

### <what>

#### *expression*

Almost any C expression, including function calls (must be prefixed with a cast to tell GDB the return value type).

#### *file\_name::variable\_name*

Content of the variable defined in the named file (static variables).

#### *function::variable\_name*

Content of the variable defined in the named function (if on the stack).

#### *{type}address*

Content at *address*, interpreted as being of the C type *type*.

#### *\$register*

Content of named register. Interesting registers are \$esp (stack pointer), \$ebp (frame pointer) and \$eip (instruction pointer).

### Threads

#### *thread <thread#>*

Chose thread to operate on.

### Manipulating the program

*set var <variable\_name>=<value>*  
Change the content of a variable to the given value.

*return <expression>*  
Force the current function to return immediately, passing the given value.

### Sources

*directory <directory>*  
Add *directory* to the list of directories that is searched for sources.

*list*  
*list <filename>:<function>*  
*list <filename>:<line\_number>*  
*list <first>,<last>*  
Shows the current or given source context. The *filename* may be omitted. If *last* is omitted the context starting at *start* is printed instead of centered around it.

*set listsize <count>*  
Set how many lines to show in „list“.

### Signals

*handle <signal> <options>*  
Set how to handle signals. Options are:

- (no)print*: (Don't) print a message when signals occurs.
- (no)stop*: (Don't) stop the program when signals occurs.
- (no)pass*: (Don't) pass the signal to the program.

### Informations

*disassemble*  
*disassemble <where>*  
Disassemble the current function or given location.

*info args*  
Print the arguments to the function of the current stack frame.

*info breakpoints*  
Print informations about the break- and watchpoints.

*info display*  
Print informations about the „displays“.

*info locals*  
Print the local variables in the currently selected stack frame.

*info sharedlibrary*  
List loaded shared libraries.

*info signals*  
List all signals and how they are currently handled.

*info threads*  
List all threads.

*show directories*  
Print all directories in which GDB searches for source files.

*show listsize*  
Print how many are shown in the „list“ command.

*whatis variable\_name*  
Print type of named variable.

# 8-Thread/OpenMP

- Race condition. Can you spot it/them ?

```
#include <omp.h>
#include <iostream>
#include <vector>

int main(int argc, char* argv[]) {
    std::vector<double> data(100, 1.0);
    double sum;

    #pragma omp parallel shared(sum, data)
    {
        sum = 0.0;
        #pragma omp for
        for(int i = 0; i < data.size(); i++){
            sum += data[i];
        }
    }
    std::cout << sum << std::endl;
}
```

→ clang++ -fopenmp -g -fsanitize=thread 8-RaceCondition.cpp

- Clang needs to be configured with -DLIBOMP\_TSAN\_SUPPORT=1

**Thank you for your attention!**